
1. Title

0:10 — open cold. Do not read the title or the tagline off the slide.

Say: *“I want to tell you about a man in Haveri who opened an envelope.”* Then advance.

Nothing on this slide needs explaining. The tagline lands at the end, not here.

2. Haveri, 2025

0:45 — the hook. Tell it, do not summarise it.

Four beats, in this order: he sells vegetables; vegetables are exempt from GST; he took UPI because *we* asked him to; the department added up four years of credits and sent him a bill for **₹29 lakh**.

Pause on the number. Then the card: within weeks, shutters across Karnataka read “No UPI, cash only.” Merchants we had just onboarded, boycotting the rail.

Close on the bottom line — the notices were withdrawn for traders who could *show* what they sold. Most could not. That verb is the whole product.

Do not say this is happening right now. It is 2025. The next slide is what makes it current, and a judge who knows the episode will respect the distinction.

3. And this is not a 2025 story

0:25 — fast. This slide exists to kill one objection.

Say: *“That was last year. This was published ten days ago.”*

Do not read the quote. Point at the two bolded phrases only — **differences in turnover** and **exemption claims** — and say those two categories are the entire content of what we build.

One extra fact if you have a beat spare: the scrutiny has moved from *unregistered* traders to registered ones, composition dealers first.

4. Every rupee is now evidence

0:40 — two authorities, one unanswerable question.

Left card briskly: the department reconciles UPI collections against filed returns and bills the difference.

Right card is the one that lands — fraud money layers fast, and the shop at the end of the chain, *the one that sold someone a watermelon*, gets frozen. Not named in the FIR, which was filed in another state. Use the watermelon.

Then the pivot: both demand the same artefact — what was each rupee, actually.

Slow down for the red line. This is the business case, not a footnote: a merchant who cannot explain their receipts takes the QR code down. Dead Soundbox, zero GMV, no lending origination, permanently.

5. Gross credits are not turnover

0:45 — the hinge of the pitch. Let the picture work.

Say: *“One shop, one year, ₹60.98 lakh of credits. The red bar is what the notice counts as turnover — all of it.”*

Then the bottom bar. **Stop talking for two seconds** and let them read it.

₹42.9L really is turnover, and we say so. ₹18.1L is not — and the point to make out loud is that **the gap is not evasion**: it is her savings account, her husband, a chit payout and a double-tap on the Soundbox.

If anyone challenges the numbers: this is seeded ground truth from our generator, not an estimate. Accuracy against it is slide 14.

6. The insight

0:40 — why this is one product, not two.

Say: *“Most teams would build the tax tool or the fraud tool. They are the same engine.”* Both are downstream reads of one ledger — two output formats.

Then the wedge, and say it slowly because it is the defensible-moat answer: Khatabook and Vyapar keep books after the fact, ClearTax files after the fact. Neither sits on the rail at the moment money moves. **Paytm does.**

Landing line: classification is cheap and accurate at the instant of the transaction, and unrecoverable eighteen months later when the notice arrives.

7. The engine

0:45 — the mechanism. Do not read the tag list.

Gesture at the eight tags, say “eight labels, every incoming credit gets one,” and move to the right-hand card, which is the real content.

The line that matters: you cannot infer “my husband sent this” from a QR payment. So the agent proposes a tag *with its reason* and she confirms with one tap. This is not a fallback — a ledger she attested contemporaneously survives a hearing; an algorithm’s guess produced under duress a year later does not.

Then the budget: at most three questions a day, and we measured 189 across a year, median 1.3. Ask more than that and the product is dead.

If asked about cost: 98.6% are ordinary sales settled by rules. The model only sees the 1.4% that need judgement.

8. What the merchant sees

0:20 — a glance slide. Do not tour the UI.

One sentence: everything classified, how close she is to the threshold, and the packs, on one surface.

Then the WhatsApp point — the same agent, where she actually is — and advance.

If someone asks whether this is built: the chat and the tools are live; this screen is the product surface we are designing toward.

9. Built on Phinite

0:40 — for the technical judges.

Two agent graphs, not four loose agents: `hisaab-watch` runs autonomously on cron, `hisaab-merchant` is the conversation. The nightly pass writes a queue; the chat reads it when she opens it — deliberately not a graph-to-graph call, because a queue that outlives either run is what makes the handoff robust.

Spend your time on the right-hand card. The strongest line in it: the Evidence agent is read-only — *the agent that builds the defence cannot edit the evidence*. And only the Merchant agent can commit an attested tag.

Those are enforced as Phinite tool policies, not conventions, and a denied call shows up in the trace.

10. Four decisions

0:35 — the slide that separates you from a demo.

Lead with #1: all arithmetic is deterministic Python, not model output. Wrapping arithmetic in an LLM is the standard tell, and the trace would show it.

Then #3: refusing to guess is a feature. An `unclassified` credit costs one tap; a confident wrong tag inside an evidence pack is a catastrophe.

Skim #2 and #4 — one clause each. #4 is worth a beat if the room is technical: we generated ground truth, so accuracy is a number instead of a claim.

11. Demo — beats 1 and 2

0:50 — or hand over to the live demo here.

Beat 1: Tuesday, 39 credits came in over the weekend, the agent asks about **three**. Read the Kannada line if you can; otherwise read the gloss under it. Note the husband paid *at the shop QR like any customer* — that is the case no rule can settle.

Beat 2 is the safety beat. **Say the red line out loud:** the agent's first real act is telling her she owes a registration. The crossing date is computed from classified turnover, not typed.

Before going on stage: confirm the attestation queue is whole — `pending_total` must be 3. A rehearsal that taps through beat 1 spends the queue, and Render needs a clear-cache redeploy to restore it.

12. Demo — beat 3, the notice

0:45 — this is the product. Slow down.

The notice claims ₹60.98 lakh. The pack answers it line by line, and every figure traces to transaction IDs and the attestation behind them — click one in the trace if the demo is live.

Then the honest half, which is what makes the whole pitch safe: **the claim is wrong, and she did cross ₹40 lakh on 14 March, and she must register.**

If a judge diffs this against slide 5: slide 5 shows seeded truth (₹42.9L), this shows what our system computed (₹42.2L). The 1.6% between them is the accuracy claim — say that, do not hedge.

13. Demo — beat 4, the emergency

0:40 — the moment they remember.

Set the stakes in one sentence: account frozen at half nine, supplier payment due, 339 payments came in that week, and she is not named in the FIR.

The agent isolates ₹4,200 by UTR *and independently* by amount and date, and brings the POS bill with it — 12kg of onions, till POS01, 19:47.

The decoy is the point. A second ₹4,200 that same week from a regular customer, listed in the pack and explicitly not chosen. Anyone can match an amount; the pack shows its work.

Finish on the red line: it never asserts she is innocent. Some frozen accounts genuinely are mule accounts. We assemble evidence; a human decides, and Escalation holds the pack for approval.

14. What we can actually prove

0:30 — do not read all four. Pick two.

Lead with **1.6%** error on turnover and **1.3 questions a day**. The other two are there to be read, not narrated.

The credibility move is the box on the right — volunteer the weakness before anyone finds it. A rules-only baseline gets 99.3% on sale vs not-a-sale but catches only 77.8% of non-sale credits, and that residue alone puts two eval merchants on the wrong side of ₹40 lakh. Closing it is what the agent plus one tap is for.

Last line if you have the beat: she corrected the agent 52 times that year, and those corrections are the ledger's strongest evidence, not its failures.

15. The line we do not cross

0:30 — the non-negotiables, fast, then land.

Run the five bullets at pace — they are a list, not five points. The one to hit is the first: we compute what she genuinely owes and say so plainly.

Then the navy box: the rail is the only place provenance is cheap. Everyone else arrives eighteen months late. And the alternative is the sign on the shutter.

Final line, unhurried: *“Paytm Hisaab. Every rupee, on your side.”* Stop talking.

— Q&A prep —

“Threshold projection but no return mismatch?” That check needs a *registered* composition dealer — someone who filed a CMP-08 to disagree with. Ours is unregistered, which is why she gets a notice at all. The composition dealer in our eval split has four quarters of under-declared returns seeded, and the same ledger answers it.

“Exempt vs taxable on a QR sale with no item list?” We do not guess per payment. We apportion by value from that shop’s own POS-billed exempt ratio and label it an estimate. Labelling each unbilled sale by its majority side understated taxable turnover threefold — the one direction we must never drift.

“Four LLM calls in a trenchcoat?” Four agents is the ceiling we set ourselves, and everything deterministic is a skill. If you can replace an agent with twenty lines of Python it should have been a skill — the trace shows which is which.

“Paytm has no standing with the police.” Agreed. Paytm supplies the evidence, the merchant files. That is the design, not a late discovery. And on cash sales: real, out of scope, and already on the slide.